

A Comparative Study of Ridge and Lasso Regularization Based on MSE Performance on Real Estate Data


Simple example to understand how L1 (Lasso) and L2 (Ridge) regularization work in practice.

Scenario: Predicting House Prices Imagine you have a dataset with the following features: Size (sqft) Bedrooms Age (years) Distance to city (km) Noise Level Price (\$)

Let's say you're using linear regression to predict house prices based on these 4 features.



Simple example to understand how L1 (Lasso) and L2 (Ridge) regularization work in practice. Scenario: Predicting House Prices Imagine you have a dataset with the following features: Size (sqft) Bedrooms Age (years) Distance to city (km) Noise Level Price (\$)

Let's say you're using linear regression to predict house prices based on these 4 features.

 Step 1: The Problem Without Regularization A linear model would look like this:

Price

$w_1 \cdot \text{Size} + w_2 \cdot \text{Bedrooms} + w_3 \cdot \text{Age} + w_4 \cdot \text{Distance} + w_5 \cdot \text{Noise} + b$
 $\text{Price} = w_1 \cdot \text{Size} + w_2 \cdot \text{Bedrooms} + w_3 \cdot \text{Age} + w_4 \cdot \text{Distance} + w_5 \cdot \text{Noise} + b$ When trained without regularization, the model might overfit — it memorizes the training data, including noise, and generalizes poorly to new data.

 Step 2: L2 Regularization (Ridge) Let's say Ridge regression learns: • $w_1 = 250$ • $w_2 = 10000$ • $w_3 = -500$ • $w_4 = -3000$ • $w_5 = -50$ Here, none of the weights are zero — Ridge shrinks them but keeps all features. Even less important features like "Noise Level" are kept with small weights.  Good when all features have some relevance.

✂ Step 3: L1 Regularization (Lasso) Lasso regression might learn: • $w_1 = 300$ • $w_2 = 0$ • $w_3 = -600$ • $w_4 = -3500$ • $w_5 = 0$ Here, some weights are exactly zero. Lasso eliminated "Bedrooms" and "Noise Level" — deciding they're not useful for predicting price. ✅ Good when you want automatic feature selection.

💡 Intuition Recap • Ridge: "Keep all features, just shrink their impact." • Lasso: "Keep only the most important features; ignore the rest."

📊 Step 2: L2 Regularization (Ridge) Let's say Ridge regression learns: • $w_1 = 250$ • $w_2 = 10000$ • $w_3 = -500$ • $w_4 = -3000$ • $w_5 = -50$ Here, none of the weights are zero — Ridge shrinks them but keeps all features. Even less important features like "Noise Level" are kept with small weights. ✅ Good when all features have some relevance.

✂ Step 3: L1 Regularization (Lasso) Lasso regression might learn: • $w_1 = 300$ • $w_2 = 0$ • $w_3 = -600$ • $w_4 = -3500$ • $w_5 = 0$ Here, some weights are exactly zero. Lasso eliminated "Bedrooms" and "Noise Level" — deciding they're not useful for predicting price. ✅ Good when you want automatic feature selection.

💡 Intuition Recap • Ridge: "Keep all features, just shrink their impact." • Lasso: "Keep only the most important features; ignore the rest."

```
In [1]: import pandas as pd
```

```
In [6]: # Load the dataset
data = pd.read_csv("customer_satisfaction_vs_repeat_purchases.csv")
```

```
In [24]: data
```

Out[24]:

	Product_Quality	Customer_Service	Delivery_Speed	Ease_of_Use	Value_for_Money
0	7	4	8	5	7
1	4	8	5	7	6
2	8	5	8	7	7
3	7	3	10	9	7
4	3	7	3	9	6
5	3	7	2	4	5
6	1	5	1	1	2
7	9	8	4	7	8
8	10	10	4	2	7
9	7	3	9	6	6
10	8	9	7	10	8
11	8	8	8	7	7
12	10	9	10	10	9
13	6	1	7	9	6
14	2	8	8	5	5
15	4	9	7	9	6
16	9	9	10	10	8
17	6	6	6	6	5
18	7	10	8	8	7
19	4	9	2	10	6
20	5	6	2	4	4
21	1	4	2	3	2
22	9	10	10	10	8
23	9	10	10	10	9
24	6	8	6	9	7
25	5	7	4	5	5
26	8	6	8	9	7
27	10	4	7	8	8
28	9	2	9	3	6
29	6	2	10	5	6
30	6	7	8	10	7
31	4	3	2	2	2
32	9	10	7	10	8

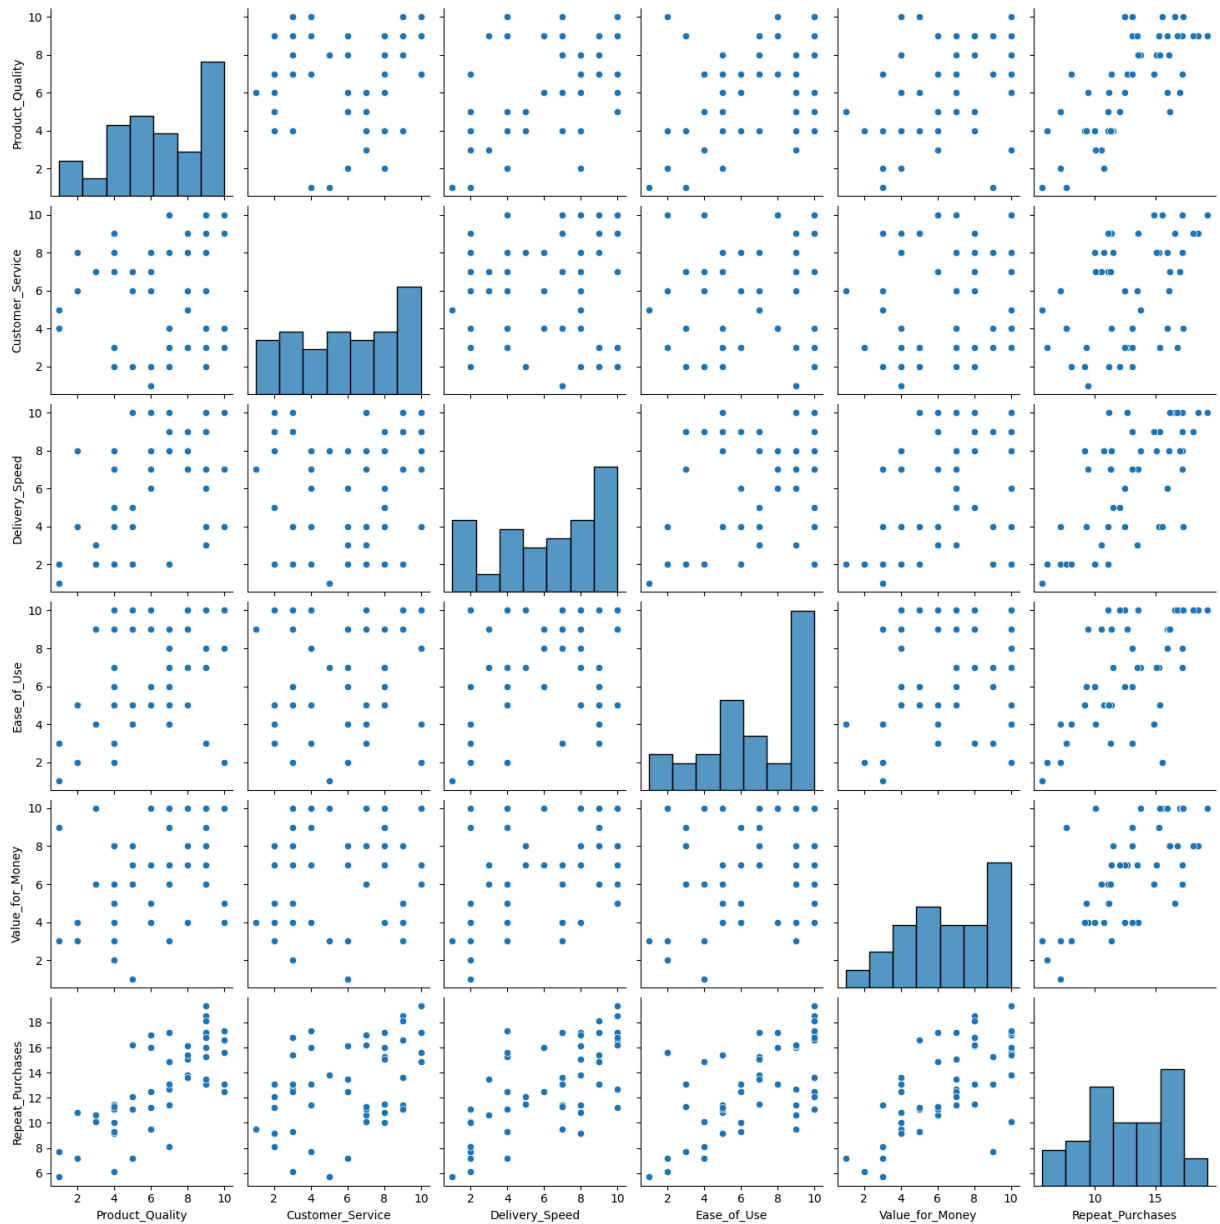
	Product_Quality	Customer_Service	Delivery_Speed	Ease_of_Use	Value_for_Money
33	9	9	9	10	
34	4	2	8	5	
35	9	4	6	8	
36	2	6	4	2	
37	9	6	3	7	
38	5	7	10	9	
39	4	8	2	6	
40	8	3	9	5	
41	10	3	4	10	
42	7	8	9	7	
43	7	2	2	4	
44	4	3	4	6	
45	9	3	10	10	
46	10	4	4	10	
47	4	7	7	3	
48	7	10	9	4	
49	5	2	5	10	

```
In [8]: data.columns
```

```
Out[8]: Index(['Product_Quality', 'Customer_Service', 'Delivery_Speed', 'Ease_of_Use',
              'Value_for_Money', 'Repeat_Purchases'],
              dtype='object')
```

```
In [9]: import seaborn as sn
        sn.pairplot(data)
```

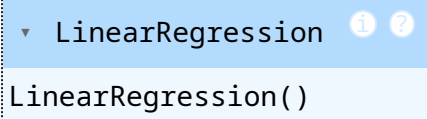
```
Out[9]: <seaborn.axisgrid.PairGrid at 0x786f64af1b50>
```



```
In [11]: # Features and target
X = data[['Customer_Service', 'Delivery_Speed', 'Ease_of_Use',
          'Value_for_Money', 'Repeat_Purchases']]
y = data["Product_Quality"]
```

```
In [12]: # Split the dataset
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, ran
```

```
In [13]: # Linear Regression (no regularization)
from sklearn.linear_model import LinearRegression
lr = LinearRegression()
lr.fit(X_train, y_train)
```

Out[13]:  LinearRegression()

```
In [14]: print("Linear Regression Coefficients:", lr.coef_)
```

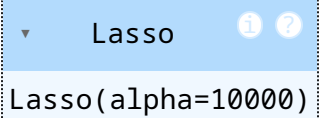
Linear Regression Coefficients: [-0.51480574 -0.32930405 -0.20035858 -0.57348732 1.29952673]

```
In [15]: from sklearn.metrics import mean_squared_error
y_pred = lr.predict(X_test)
mse_No_regulation = mean_squared_error(y_test, y_pred)
mse_No_regulation
```

Out[15]: 1.2194994717157444

```
In [16]: # Linear Regression (L1 regularization)
from sklearn.linear_model import Lasso
lasso = Lasso(alpha=10000)

lasso.fit(X_train, y_train)
```

Out[16]:  Lasso(alpha=10000)

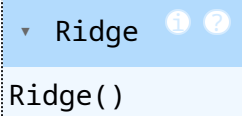
```
In [17]: print("Lasso Coefficients:", lasso.coef_)
```

Lasso Coefficients: [0. 0. 0. 0. 0.]

```
In [18]: from sklearn.metrics import mean_squared_error
y_pred = lasso.predict(X_test)
mse_L1_regulation = mean_squared_error(y_test, y_pred)
mse_L1_regulation
```

Out[18]: 5.119727891156462

```
In [19]: # Linear Regression (L2 regularization)
from sklearn.linear_model import Ridge
ridge = Ridge(alpha=1.0)
ridge.fit(X_train, y_train)
```

Out[19]:  Ridge()

```
In [20]: print("Lasso Coefficients:", ridge.coef_)
```

Lasso Coefficients: [-0.50487409 -0.3112752 -0.17493312 -0.54869774 1.25637873]

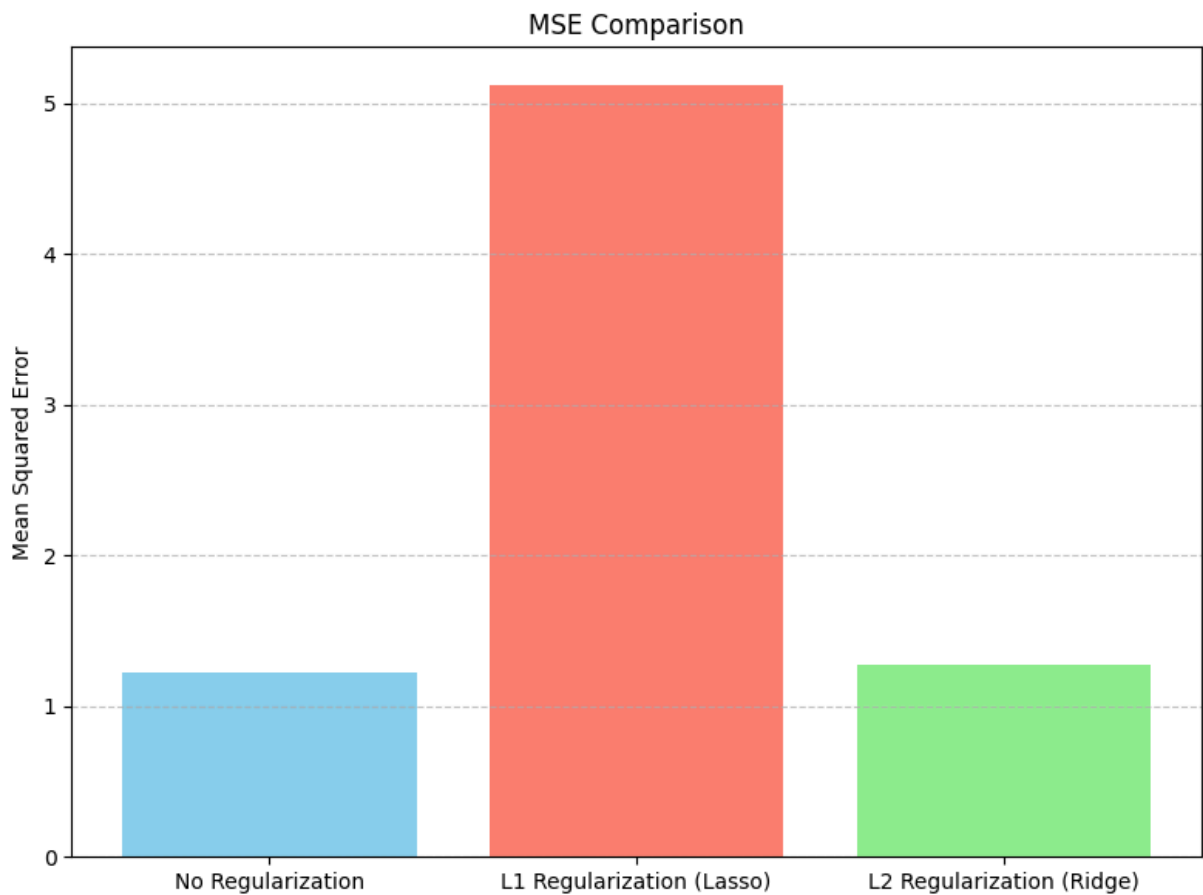
```
In [21]: from sklearn.metrics import mean_squared_error
y_pred = ridge.predict(X_test)
mse_L2_regulation = mean_squared_error(y_test, y_pred)
mse_L2_regulation
```

Out[21]: 1.2773266643700885

```
In [22]: import matplotlib.pyplot as plt

mse_values = {
    "No Regularization": mse_No_regulation,
    "L1 Regularization (Lasso)": mse_L1_regulation,
    "L2 Regularization (Ridge)": mse_L2_regulation
}

plt.figure(figsize=(8, 6))
plt.bar(mse_values.keys(), mse_values.values(), color=["skyblue", "salmon",
plt.title("MSE Comparison")
plt.ylabel("Mean Squared Error")
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```



```
In [23]: # no regulation and l2 is same because the data is cleaned
```